


Cross-system comparison on the DDI test set — small quantised LLMs trail dedicated NERC systems
(macro gap 13 pp, micro gap 28 pp — the LLM under-performs most on the common classes)

Test micro-F1 (m.avg)



Test macro-F1 (M.avg)

